

Mykhailo Kapustin

kapustinomm@gmail.com · Berlin, Germany · github.com/Kapustin2000/ · linkedin.com/in/mykhailo-kapustin

Dear InsulaLabs team,

I'm writing about the Senior Data Engineer role. Identity resolution with deep AI/LLM integration in production pipelines at 10,000+ operations a minute is a specific, hard problem, and it's conceptually the same class of work I've spent time on.

Why this role fits. Docling in Working with Texts, Languages, and Knowledge covers graph-based knowledge representation — parent-child and semantic links between entities in a document — which is the same underlying logic as identity resolution applied to records about people or entities, just a different data domain. As the lead engineer (through my company, as the contracting entity) on a multi-tenant agent platform for a German AI SaaS company, I built the Python/FastAPI backend for dynamic agents and production data pipelines, with PostgreSQL row-level security for tenant isolation at scale, and was the top contributor (~45%) across both core repositories.

Practical note. I hold EU citizenship and am based in Berlin. I'm open to full-time, contract, or freelance arrangements — whatever fits how you're structuring the hire.

I'd welcome the chance to learn more about the end client and industry behind this identity-resolution system, and whether this specific role leans more toward data engineering or data analysis.

Best regards,

Mykhailo Kapustin

Berlin, Germany

github.com/Kapustin2000/

kapustinomm@gmail.com