Mykhailo Kapustin

kapustinomm@gmail.com · Berlin, Germany · github.com/Kapustin2000/ · linkedin.com/in/mykhailo-kapustin

Dear Nova Digital team,

I'm writing about the Python AI Lead role. An LLM/agents platform running at 50M+ requests a day is a different kind of engineering problem than most "AI agent" work I see advertised — at that scale, agents can't be allowed to be fragile, and the operational discipline around them matters as much as the model layer itself. That's the part of the role I find most interesting.

Why this role fits. As the lead engineer (through ASRP, as the contracting entity) on a multi-tenant agent platform for a German AI SaaS company, I owned the architecture end to end: a NestJS/React/PostgreSQL customer-facing layer with row-level security for tenant isolation, natural-language agent authoring (describe an agent in plain text, get a working, deployed one back in about a minute), and the full staging-to-production release process, as the top contributor (~45%) across both core repositories. Before that, I held a formal Team Lead title (MOB.325/LAB325, Ukraine) running Node.js/GraphQL backend work. I've written about exactly this class of problem — Building AI Agent Infrastructure (10k+ reads) walks through the path from ad-hoc tool wiring to an AgentOps layer that behaves like a proper microservice system as usage scales, and Google ADK and "Startup Technical Guide: AI Agents" covers why most agents actually fail in production — missing operational discipline, not bad models.

One thing worth clarifying up front. "Lead" can mean different things — architectural ownership without direct reports, or people management. My background covers the former solidly (plus one formal Team Lead title); if the role leans heavily toward people management, I'd want to talk through what that looks like day to day before assuming I'm the right fit.

Practical note. I hold EU citizenship and am based in Berlin, and I'm open to full-time, contract, or freelance arrangements — whatever fits how you're structuring the hire.

I'd welcome the chance to talk through how the LLM/agent layer fits into the wider 50M+ requests/day infrastructure, and how much of it is agent-based today versus classic ML/rule-based.

Best regards,
Mykhailo Kapustin
Berlin, Germany
github.com/Kapustin2000/
kapustinomm@gmail.com