

Mykhailo Kapustin

kapustinomm@gmail.com · Berlin, Germany · github.com/Kapustin2000/ · linkedin.com/in/mykhailo-kapustin

Dear Workeron.ai team,

I'm writing about the Senior AI/LLM Python developer (LangGraph) role. "Configure an agentic digital employee once, deploy it everywhere" is close to a feature I've already shipped, so I wanted to reach out directly.

Why this role fits. As the lead engineer (through my company, as the contracting entity) on a multi-tenant agent platform for a German AI SaaS company, I built the Python/FastAPI backend for dynamic agents — tool registries, tenant-safe execution — with natural-language agent authoring: describe an agent in plain text, get a working, deployed agent back in ~1 minute, running across web chat, email, and Outlook. I was the top contributor (~45%) across both core repositories and owned the release process from staging to production. LangGraph has been my production framework for close to 3 years, most recently for a forecasting platform I built at an AI hackathon in Berlin and have kept iterating on since.

Why I think about this the way you do. My Personal Exam: MVP LLM Agent on Google ADK covers multi-step agent coordination and reviewer-agent validation before an answer goes out — relevant if a "digital employee" is meant to be trusted in production, not just impressive in a demo. Building AI Agent Infrastructure covers the tool-registry/tenant-safe-execution architecture that keeps agents reliable across many clients at once.

Practical note. I hold EU citizenship and am based in Berlin. I'm open to full-time, contract, or freelance arrangements — whatever fits how you're structuring the hire.

I'd welcome the chance to talk through the architecture in more depth, and to hear more about the team and stage of the product.

Best regards,

Mykhailo Kapustin

Berlin, Germany

github.com/Kapustin2000/

kapustinomm@gmail.com